

# Class\_09\_Mini\_Project

Hayoung A15531571

10/26/2021

To get started let's read the data!

```
#we tell R what file we want here
fna.data <- "WisconsinCancer.csv"

#making sure to format the data correctly
wisc.df <- read.csv(fna.data, row.names=1)

#finally let's view our data
head(wisc.df)
```

```
##      diagnosis radius_mean texture_mean perimeter_mean area_mean
## 842302         M      17.99        10.38         122.80      1001.0
## 842517         M      20.57        17.77         132.90      1326.0
## 84300903        M      19.69        21.25         130.00      1203.0
## 84348301         M      11.42        20.38          77.58       386.1
## 84358402         M      20.29        14.34         135.10      1297.0
## 843786          M      12.45        15.70          82.57       477.1
##      smoothness_mean compactness_mean concavity_mean concave.points_mean
## 842302          0.11840          0.27760          0.3001          0.14710
## 842517          0.08474          0.07864          0.0869          0.07017
## 84300903         0.10960          0.15990          0.1974          0.12790
## 84348301         0.14250          0.28390          0.2414          0.10520
## 84358402         0.10030          0.13280          0.1980          0.10430
## 843786          0.12780          0.17000          0.1578          0.08089
##      symmetry_mean fractal_dimension_mean radius_se texture_se perimeter_se
## 842302          0.2419          0.07871          1.0950          0.9053          8.589
## 842517          0.1812          0.05667          0.5435          0.7339          3.398
## 84300903         0.2069          0.05999          0.7456          0.7869          4.585
## 84348301         0.2597          0.09744          0.4956          1.1560          3.445
## 84358402         0.1809          0.05883          0.7572          0.7813          5.438
## 843786          0.2087          0.07613          0.3345          0.8902          2.217
##      area_se smoothness_se compactness_se concavity_se concave.points_se
## 842302      153.40      0.006399      0.04904      0.05373      0.01587
## 842517       74.08      0.005225      0.01308      0.01860      0.01340
## 84300903      94.03      0.006150      0.04006      0.03832      0.02058
## 84348301      27.23      0.009110      0.07458      0.05661      0.01867
## 84358402      94.44      0.011490      0.02461      0.05688      0.01885
## 843786      27.19      0.007510      0.03345      0.03672      0.01137
##      symmetry_se fractal_dimension_se radius_worst texture_worst
## 842302      0.03003      0.006193      25.38      17.33
```

```
## 842517      0.01389      0.003532      24.99      23.41
## 84300903    0.02250      0.004571      23.57      25.53
## 84348301    0.05963      0.009208      14.91      26.50
## 84358402    0.01756      0.005115      22.54      16.67
## 843786      0.02165      0.005082      15.47      23.75
##           perimeter_worst area_worst smoothness_worst compactness_worst
## 842302           184.60    2019.0           0.1622           0.6656
## 842517           158.80    1956.0           0.1238           0.1866
## 84300903         152.50    1709.0           0.1444           0.4245
## 84348301           98.87     567.7           0.2098           0.8663
## 84358402         152.20    1575.0           0.1374           0.2050
## 843786          103.40     741.6           0.1791           0.5249
##           concavity_worst concave.points_worst symmetry_worst
## 842302           0.7119           0.2654           0.4601
## 842517           0.2416           0.1860           0.2750
## 84300903         0.4504           0.2430           0.3613
## 84348301         0.6869           0.2575           0.6638
## 84358402         0.4000           0.1625           0.2364
## 843786           0.5355           0.1741           0.3985
##           fractal_dimension_worst
## 842302           0.11890
## 842517           0.08902
## 84300903         0.08758
## 84348301         0.17300
## 84358402         0.07678
## 843786           0.12440
```

When we look at our data so far, we realize that we don't want the first row which tells us right away if something is malignant or benign.

```
# We can use -1 here to remove the first column
wisc.data <- wisc.df[,-1]
```

```
# Create diagnosis vector for later
diagnosis <- as.factor(wisc.df$diagnosis)
```

Let's move onto the questions!

Q1. How many observations are in this dataset?

```
#here we read all the data (minus the first diagnosis column)
dim(wisc.data)
```

```
## [1] 569 30
```

There are 569 rows (or different biopsies to analyze). Each biopsy has 30 elements to it (30 rows).

Q2. How many of the observations have a malignant diagnosis?

```
#here we can use our diagnosis vector to see how many malignant, or "M" results we have
length(grep(pattern = "M", x = diagnosis))
```

```
## [1] 212
```

There are 212 malignant results in this data set (out of 569 biopsies)

Q3. How many variables/features in the data are suffixed with `_mean`?

```
#first we have to be able to read the column names
features <- colnames(wisc.df)
length(grep(pattern = "_mean", x = features))
```

```
## [1] 10
```

There are 10 variables with “`_mean`” in the variable name.

## Performing PCA

```
# Check column means and standard deviations
colMeans(wisc.data)
```

```
##           radius_mean      texture_mean      perimeter_mean
##      1.412729e+01      1.928965e+01      9.196903e+01
##           area_mean      smoothness_mean      compactness_mean
##      6.548891e+02      9.636028e-02      1.043410e-01
##      concavity_mean      concave.points_mean      symmetry_mean
##      8.879932e-02      4.891915e-02      1.811619e-01
## fractal_dimension_mean      radius_se      texture_se
##      6.279761e-02      4.051721e-01      1.216853e+00
##      perimeter_se      area_se      smoothness_se
##      2.866059e+00      4.033708e+01      7.040979e-03
##      compactness_se      concavity_se      concave.points_se
##      2.547814e-02      3.189372e-02      1.179614e-02
##      symmetry_se      fractal_dimension_se      radius_worst
##      2.054230e-02      3.794904e-03      1.626919e+01
##      texture_worst      perimeter_worst      area_worst
##      2.567722e+01      1.072612e+02      8.805831e+02
##      smoothness_worst      compactness_worst      concavity_worst
##      1.323686e-01      2.542650e-01      2.721885e-01
##      concave.points_worst      symmetry_worst      fractal_dimension_worst
##      1.146062e-01      2.900756e-01      8.394582e-02
```

```
apply(wisc.data,2,sd)
```

```
##           radius_mean      texture_mean      perimeter_mean
##      3.524049e+00      4.301036e+00      2.429898e+01
##           area_mean      smoothness_mean      compactness_mean
##      3.519141e+02      1.406413e-02      5.281276e-02
##      concavity_mean      concave.points_mean      symmetry_mean
##      7.971981e-02      3.880284e-02      2.741428e-02
## fractal_dimension_mean      radius_se      texture_se
```

```
##          7.060363e-03          2.773127e-01          5.516484e-01
##          perimeter_se          area_se          smoothness_se
##          2.021855e+00          4.549101e+01          3.002518e-03
##          compactness_se          concavity_se          concave.points_se
##          1.790818e-02          3.018606e-02          6.170285e-03
##          symmetry_se          fractal_dimension_se          radius_worst
##          8.266372e-03          2.646071e-03          4.833242e+00
##          texture_worst          perimeter_worst          area_worst
##          6.146258e+00          3.360254e+01          5.693570e+02
##          smoothness_worst          compactness_worst          concavity_worst
##          2.283243e-02          1.573365e-01          2.086243e-01
##          concave.points_worst          symmetry_worst          fractal_dimension_worst
##          6.573234e-02          6.186747e-02          1.806127e-02
```

Let's execute PCA now

```
# Perform PCA on wisc.data by completing the following code
wisc.pr <- prcomp(wisc.data, scale = TRUE)

# Look at summary of results
summary(wisc.pr)
```

```
## Importance of components:
##          PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  3.6444 2.3857 1.67867 1.40735 1.28403 1.09880 0.82172
## Proportion of Variance 0.4427 0.1897 0.09393 0.06602 0.05496 0.04025 0.02251
## Cumulative Proportion 0.4427 0.6324 0.72636 0.79239 0.84734 0.88759 0.91010
##          PC8      PC9      PC10     PC11     PC12     PC13     PC14
## Standard deviation  0.69037 0.6457 0.59219 0.5421 0.51104 0.49128 0.39624
## Proportion of Variance 0.01589 0.0139 0.01169 0.0098 0.00871 0.00805 0.00523
## Cumulative Proportion 0.92598 0.9399 0.95157 0.9614 0.97007 0.97812 0.98335
##          PC15     PC16     PC17     PC18     PC19     PC20     PC21
## Standard deviation  0.30681 0.28260 0.24372 0.22939 0.22244 0.17652 0.1731
## Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
## Cumulative Proportion 0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
##          PC22     PC23     PC24     PC25     PC26     PC27     PC28
## Standard deviation  0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
## Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
## Cumulative Proportion 0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
##          PC29     PC30
## Standard deviation  0.02736 0.01153
## Proportion of Variance 0.00002 0.00000
## Cumulative Proportion 1.00000 1.00000
```

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

From our summary above we can see that 44.27% of the variance is captured by PC1

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?


To describe at least 70% of variance, we need 3 principal components

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

To describe at least 90% of variance, we need 7 principal components

#Now let's try plotting this out

```
biplot(wisc.pr)
```



```
as.factor(diagnosis)
```

```
## [1] M M M M M M M M M M M M M M M M M M M M M M M M M M M M M
## [38] B M M M M M M M M M M M M M M M M M M M M M M M M M M M M
## [75] B M B M M B B B M M M M M M M M M M M M M M M M M M M M M
## [112] B B B B B B M M M M M M M M M M M M M M M M M M M M M M M
## [149] B B B B B B B B M M M M M M M M M M M M M M M M M M M M M
## [186] B M B B B M B B M M M M M M M M M M M M M M M M M M M M M
## [223] B M B B B B B M M M M M M M M M M M M M M M M M M M M M M
## [260] M M M M M M M B B B B B M M M M M M M M M M M M M M M M M
## [297] B M B B M B M B B B B B B B B B B B B B B B B M M M M M M
## [334] B B M B M B M B B B M M M M M M M M M M M M M M M M M M M
## [371] M B M M B B B B B M B B B B B M M M M M M M M M M M M M
## [408] B M B B B B B M B B M M M M M M M M M M M M M M M M M M M
## [445] M B M B B M B M B B B B B B M M M M M M M M M M M M M
## [482] B B B B B B M M M M M M M M M M M M M M M M M M M M M M
```


```
## [519] B B B M B B B B B B B B B B M B M M B B B B B B B B B B B B B B
## [556] B B B B B B B M M M M M M B
## Levels: B M
```

Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

It is very messy and unable to be read. Even when we pop it out in the larger browser, there are so many points that it is impossible to really read or understand. Row names are being used as labels which makes it hard to read, considering how many rows we have.


Let's try this again, we are after the score plot (ex: PC1 vs PC2)

```
# Scatter plot observations by components 1 and 2
plot(wisc.pr$x[,1:2], col = diagnosis ,
     xlab = "PC1", ylab = "PC2")
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
# Scatter plot observations by components 1 and 3
plot(wisc.pr$x[,1], wisc.pr$x[,3], col = diagnosis ,
     xlab = "PC1", ylab = "PC3")
```

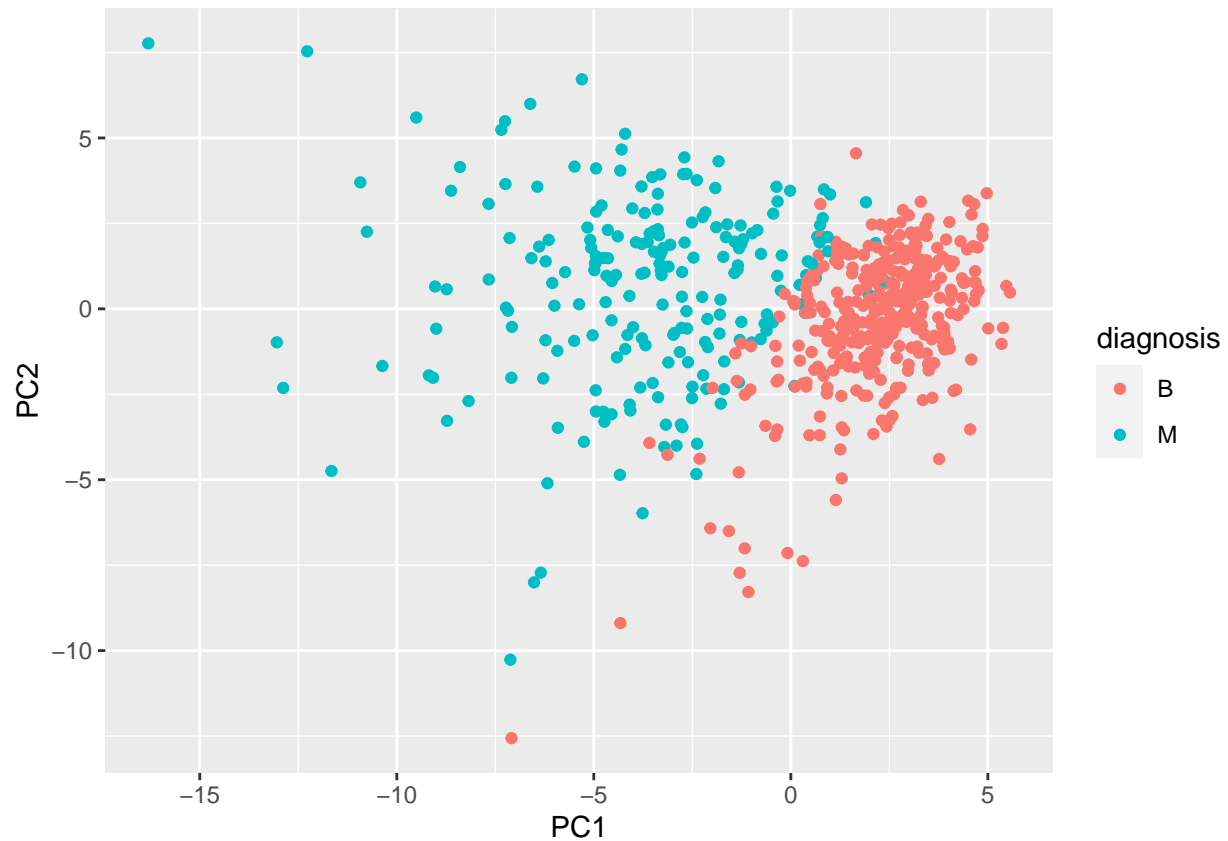


#Now let's try using ggplot to make a nicer looking plot

```
# Create a data.frame for ggplot
df <- as.data.frame(wisc.pr$x)
df$diagnosis <- diagnosis

# Load the ggplot2 package
library(ggplot2)

# Make a scatter plot colored by diagnosis
ggplot(df) +
  aes(PC1, PC2, col = diagnosis) +
  geom_point()
```



#Variance

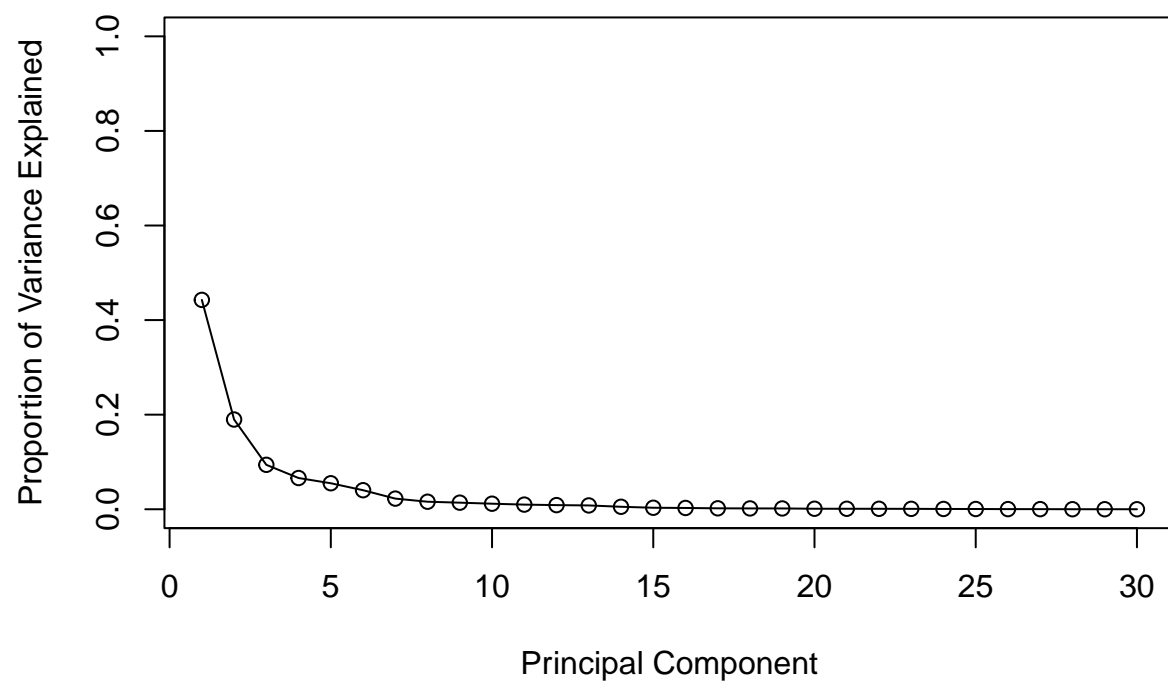
```
# Calculate variance of each component
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

```
## [1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```

```
# Variance explained by each principal component: pve
pve <- pr.var / sum(pr.var)

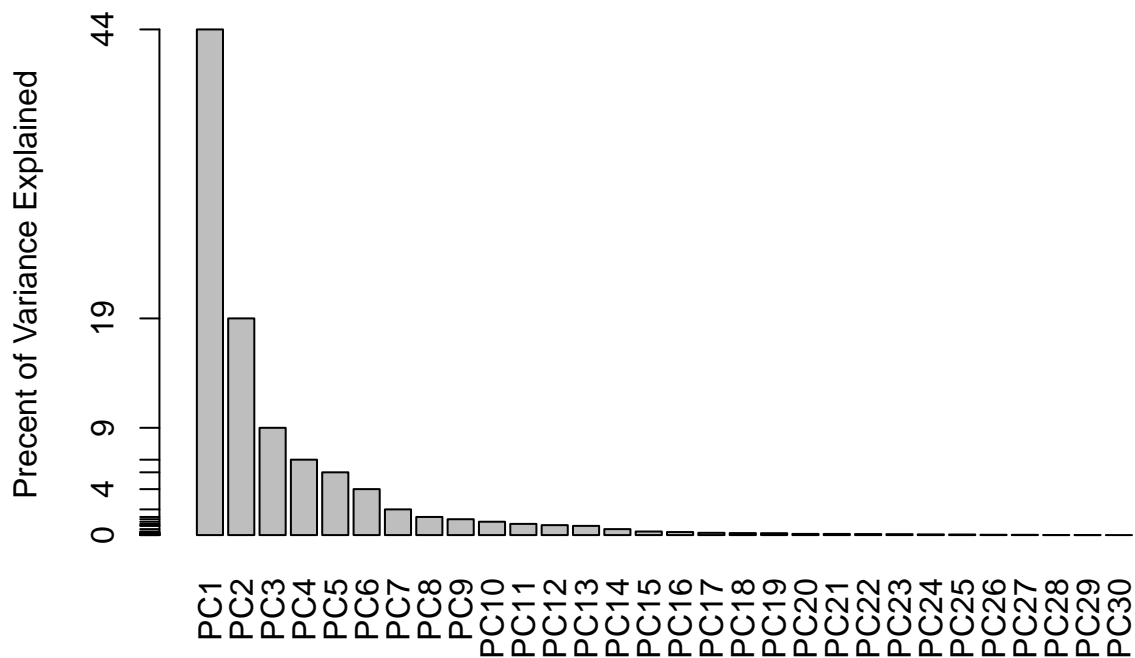
# Plot variance explained for each principal component
plot(pve, xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0, 1), type = "o")
```





An alternative graph!

```
# data driven y-axis
barplot(pve, ylab = "Precent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```



Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`?

```
wisc.pr$rotation["concave.points_mean",1]
```

```
## [1] -0.2608538
```

The component of the loading vector for the feature `concave.points_mean` is -0.26085376.

Q10. What is the minimum number of principal components required to explain 80% of the variance of the data?

```
var <- summary(wisc.pr)
var$importance[2,]
```

```
##      PC1      PC2      PC3      PC4      PC5      PC6      PC7      PC8      PC9      PC10
## 0.44272 0.18971 0.09393 0.06602 0.05496 0.04025 0.02251 0.01589 0.01390 0.01169
##      PC11     PC12     PC13     PC14     PC15     PC16     PC17     PC18     PC19     PC20
## 0.00980 0.00871 0.00805 0.00523 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104
##      PC21     PC22     PC23     PC24     PC25     PC26     PC27     PC28     PC29     PC30
## 0.00100 0.00091 0.00081 0.00060 0.00052 0.00027 0.00023 0.00005 0.00002 0.00000
```

We need at least 5 principal components to explain 80% variance of the data.

```
#Hierarchal Clustering
```

```
# Scale the wisc.data data using the "scale()" function
data.scaled <- scale(wisc.data)
```


```
#Calculate the (Euclidean) distances between all pairs of observations
data.dist <- dist(data.scaled)
```

```
#Create a hierarchical clustering model using complete linkage.
wisc.hclust <- hclust(data.dist, method = "complete")
```

Q11. Using the plot() and abline() functions, what is the height at which the clustering model has 4 clusters?

```
plot(wisc.hclust)
#not too sure how exactly to calculate where it cuts off, eyeballed it for now
abline(h = 19.5, col="red", lty=2)
```

## Cluster Dendrogram



```
data.dist
hclust (*, "complete")
```

```
#Selecting number of clusters
```

```
wisc.hclust.clusters <- cutree(wisc.hclust, k = 4)
```

```
#We can use the table() function to compare the cluster membership to the actual diagnoses.
table(wisc.hclust.clusters, diagnosis)
```

```
##          diagnosis
```

```
## wisc.hclust.clusters  B  M
##                      1 12 165
##                      2  2  5
##                      3 343 40
##                      4  0  2
```

Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?

```
#Let's try out different cluster groups
wisc.hclust.clusters2 <- cutree(wisc.hclust, k = 2)
table(wisc.hclust.clusters2, diagnosis)
```

```
##                      diagnosis
## wisc.hclust.clusters2  B  M
##                      1 357 210
##                      2  0  2
```

```
wisc.hclust.clusters3 <- cutree(wisc.hclust, k = 3)
table(wisc.hclust.clusters3, diagnosis)
```

```
##                      diagnosis
## wisc.hclust.clusters3  B  M
##                      1 355 205
##                      2  2  5
##                      3  0  2
```

```
wisc.hclust.clusters5 <- cutree(wisc.hclust, k = 5)
table(wisc.hclust.clusters5, diagnosis)
```

```
##                      diagnosis
## wisc.hclust.clusters5  B  M
##                      1 12 165
##                      2  0  5
##                      3 343 40
##                      4  2  0
##                      5  0  2
```

```
wisc.hclust.clusters6 <- cutree(wisc.hclust, k = 6)
table(wisc.hclust.clusters6, diagnosis)
```

```
##                      diagnosis
## wisc.hclust.clusters6  B  M
##                      1 12 165
##                      2  0  5
##                      3 331 39
##                      4  2  0
##                      5 12  1
##                      6  0  2
```

```
wisc.hclust.clusters7 <- cutree(wisc.hclust, k = 7)
table(wisc.hclust.clusters7, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters7  B  M
##              1 12 165
##              2  0  3
##              3 331  39
##              4  2  0
##              5 12  1
##              6  0  2
##              7  0  2
```

```
wisc.hclust.clusters8 <- cutree(wisc.hclust, k = 8)
table(wisc.hclust.clusters8, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters8  B  M
##              1 12  86
##              2  0  79
##              3  0  3
##              4 331  39
##              5  2  0
##              6 12  1
##              7  0  2
##              8  0  2
```

```
wisc.hclust.clusters9 <- cutree(wisc.hclust, k = 9)
table(wisc.hclust.clusters9, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters9  B  M
##              1 12  86
##              2  0  79
##              3  0  3
##              4 331  39
##              5  2  0
##              6 12  0
##              7  0  2
##              8  0  2
##              9  0  1
```

```
wisc.hclust.clusters10 <- cutree(wisc.hclust, k = 10)
table(wisc.hclust.clusters10, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters10  B  M
##              1 12  86
##              2  0  59
##              3  0  3
##              4 331  39
```

```
##           5    0  20
##           6    2   0
##           7   12   0
##           8    0   2
##           9    0   2
##          10    0   1
```

With 6-10 clusters, the results are almost exactly the same to one another, we can rule those out.

```
wisc.hclust.clusters2 <- cutree(wisc.hclust, k = 2)
table(wisc.hclust.clusters2, diagnosis)
```

```
##           diagnosis
## wisc.hclust.clusters2  B  M
##           1 357 210
##           2   0   2
```

```
wisc.hclust.clusters3 <- cutree(wisc.hclust, k = 3)
table(wisc.hclust.clusters3, diagnosis)
```

```
##           diagnosis
## wisc.hclust.clusters3  B  M
##           1 355 205
##           2   2   5
##           3   0   2
```

```
#this is our original
wisc.hclust.clusters <- cutree(wisc.hclust, k = 4)
table(wisc.hclust.clusters, diagnosis)
```

```
##           diagnosis
## wisc.hclust.clusters  B  M
##           1  12 165
##           2   2   5
##           3 343  40
##           4   0   2
```


```
wisc.hclust.clusters5 <- cutree(wisc.hclust, k = 5)
table(wisc.hclust.clusters5, diagnosis)
```

```
##           diagnosis
## wisc.hclust.clusters5  B  M
##           1  12 165
##           2   0   5
##           3 343  40
##           4   2   0
##           5   0   2
```

Having 3 clusters gives a better cluster vs. diagnosis match because we can see in this table that cluster 1 represents almost all the data (560/569). We can see here that there are 355 benign cells and 205 malignant ones.

Q13. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning.

```
wisc.hclust13 <- hclust(data.dist, method = "ward.D2")
plot(wisc.hclust13)
```



```
data.dist
hclust (*, "ward.D2")
```

This gives 3 clusters, like we decided on in the previous question. This gives us the grouping that we want.

#Combining Methods

We take the results of our PCA analysis and cluster in this space

```
summary(wisc.pr)
```

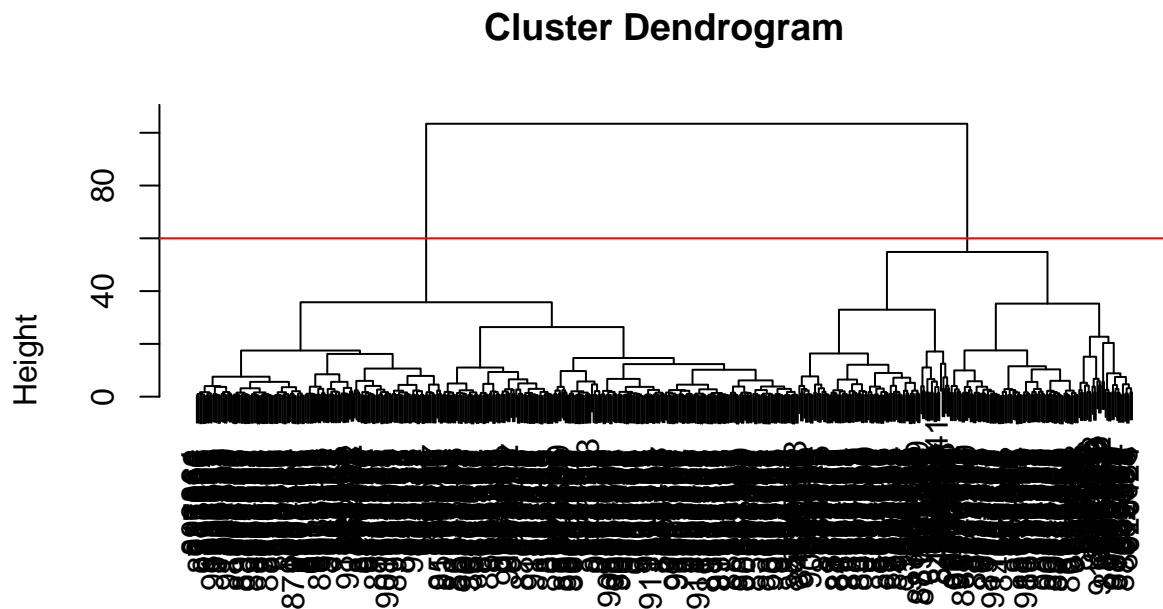
```
## Importance of components:
##              PC1    PC2    PC3    PC4    PC5    PC6    PC7
## Standard deviation  3.6444 2.3857 1.67867 1.40735 1.28403 1.09880 0.82172
## Proportion of Variance 0.4427 0.1897 0.09393 0.06602 0.05496 0.04025 0.02251
## Cumulative Proportion 0.4427 0.6324 0.72636 0.79239 0.84734 0.88759 0.91010
##              PC8    PC9    PC10   PC11   PC12   PC13   PC14
## Standard deviation  0.69037 0.6457 0.59219 0.5421 0.51104 0.49128 0.39624
## Proportion of Variance 0.01589 0.0139 0.01169 0.0098 0.00871 0.00805 0.00523
## Cumulative Proportion 0.92598 0.9399 0.95157 0.9614 0.97007 0.97812 0.98335
##              PC15   PC16   PC17   PC18   PC19   PC20   PC21
## Standard deviation  0.30681 0.28260 0.24372 0.22939 0.22244 0.17652 0.1731
## Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
```

```
## Cumulative Proportion 0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
##                      PC22  PC23  PC24  PC25  PC26  PC27  PC28
## Standard deviation    0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
## Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
## Cumulative Proportion 0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
##                      PC29  PC30
## Standard deviation    0.02736 0.01153
## Proportion of Variance 0.00002 0.00000
## Cumulative Proportion 1.00000 1.00000
```

```
#First we have to create something with at least 90% of variance explained for
wisc.pr.hclust <- hclust(dist(wisc.pr$x[,1:3] ),
                        method = "ward.D2")
```

Plotting the dendrogram

```
plot(wisc.pr.hclust)
abline(h=60, col="red")
```



```
dist(wisc.pr$x[, 1:3])
hclust (*, "ward.D2")
```

Cut the tree into k=2 groups

```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
## grps
```




```
##    1    2
## 203 366
```

Cross table compare of diagnosis and my cluster groups

```
table(diagnosis, grps)
```

```
##           grps
## diagnosis    1    2
##           B   24 333
##           M  179  33
```

```
plot(wisc.pr$x[,1:2], col=diagnosis)
```




```
g <- as.factor(grps)
levels(g)
```

```
## [1] "1" "2"
```

```
g <- relevel(g,2)
levels(g)
```

```
## [1] "2" "1"
```

```
# Plot using our re-ordered factor
plot(wisc.pr$x[,1:2], col=g)
```



```
## Use the distance along the first 7 PCs for clustering i.e. wisc.pr$x[, 1:7]
wisc.pr.hclust <- hclust(dist(wisc.pr$x[,1:7] ),
                        method = "ward.D2")
```

```
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)
```

Q15. How well does the newly created model with four clusters separate out the two diagnoses?

```
# Compare to actual diagnoses
table(wisc.pr.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.pr.hclust.clusters  B  M
##              1  28 188
##              2 329  24
```

The newly created model sorts once again benign/malignant tumors pretty well. A little better than our previous sorting but honestly still pretty similar.

Q16. How well do the k-means and hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the table() function to

compare the output of each model (wisc.km\$cluster and wisc.hclust.clusters) with the vector containing the actual diagnoses.

```
#this was from the optional part 4 but have to create a wisc.km  
wisc.km <- kmeans(wisc.data, centers= 2, nstart= 20)
```

```
table(wisc.km$cluster, diagnosis)
```

```
##      diagnosis  
##      B      M  
##  1    1  130  
##  2  356   82
```

```
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis  
## wisc.hclust.clusters  B      M  
##              1    12  165  
##              2     2    5  
##              3   343   40  
##              4     0    2
```

From our wisc.hclust.clusters (hierarchal) data, we can see that there are more clusters (4), and provide mostly better results of malignant vs benign. Cluster 1 has 12/165 showing benign, cluster 3 has 343/383 showing benign, and cluster 4 has 2/2 showing malignant. These results could lead to less false positives. But cluster 3 shows 2/7 as benign, this one cluster is not as defined (in terms of malignant vs. benign).

For our k-means, we have 2 clusters with pretty good separation of benign vs. malignant. Cluster 1 has 130/131 showing to be malignant, and Cluster 2 shows 356/438 to be benign. Personally I think the k-means data is better to look at just as a table, it's Cluster 1 is very accurate.

#Sensitivity/Specificity

Q17. Which of your analysis procedures resulted in a clustering model with the best specificity?  
How about sensitivity?

```
#Calculating for sensitivity  
130 / (130+82)
```

```
## [1] 0.6132075
```

```
165 / (165+40+5+2)
```

```
## [1] 0.7783019
```

```
#Calculating for Specificity  
356 / (356+1)
```

```
## [1] 0.9971989
```

343 / (343+12)

```
## [1] 0.9661972
```

The use of hierarchal clustering is better for sensitivity


The use of k-means is better for specificity

#Prediction

```
#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

```
##          PC1          PC2          PC3          PC4          PC5          PC6          PC7
## [1,]  2.576616 -3.135913  1.3990492 -0.7631950  2.781648 -0.8150185 -0.3959098
## [2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
##          PC8          PC9          PC10          PC11          PC12          PC13          PC14
## [1,] -0.2307350 0.1029569 -0.9272861 0.3411457  0.375921 0.1610764 1.187882
## [2,] -0.3307423 0.5281896 -0.4855301 0.7173233 -1.185917 0.5893856 0.303029
##          PC15          PC16          PC17          PC18          PC19          PC20
## [1,] 0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
## [2,] 0.1299153 0.1448061 -0.40509706 0.06565549 0.25591230 -0.4289500
##          PC21          PC22          PC23          PC24          PC25          PC26
## [1,] 0.1228233 0.09358453 0.08347651 0.1223396 0.02124121 0.078884581
## [2,] -0.1224776 0.01732146 0.06316631 -0.2338618 -0.20755948 -0.009833238
##          PC27          PC28          PC29          PC30
## [1,] 0.220199544 -0.02946023 -0.015620933 0.005269029
## [2,] -0.001134152 0.09638361 0.002795349 -0.019015820
```

```
plot(wisc.pr$x[,1:2], col=g)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q18. Which of these new patients should we prioritize for follow up based on your results?

We would prioritize patient 2 because their clustering/data looks more like the (mostly) malignant patients (more like cluster 1 which had mostly malignant tumors in our previous Wisconsin data).

```
sessionInfo()
```

```
## R version 4.1.1 (2021-08-10)
## Platform: x86_64-apple-darwin17.0 (64-bit)
## Running under: macOS Big Sur 10.16
##
## Matrix products: default
## BLAS:   /Library/Frameworks/R.framework/Versions/4.1/Resources/lib/libRblas.0.dylib
## LAPACK: /Library/Frameworks/R.framework/Versions/4.1/Resources/lib/libRlapack.dylib
##
## locale:
## [1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8
##
## attached base packages:
## [1] stats      graphics  grDevices  utils      datasets  methods   base
##
## other attached packages:
## [1] ggplot2_3.3.5
##
## loaded via a namespace (and not attached):
```

```
## [1] knitr_1.36      magrittr_2.0.1  tidyselect_1.1.1 munsell_0.5.0
## [5] colorspace_2.0-2 R6_2.5.1       rlang_0.4.11    fastmap_1.1.0
## [9] fansi_0.5.0     dplyr_1.0.7    stringr_1.4.0   highr_0.9
## [13] tools_4.1.1     grid_4.1.1     gtable_0.3.0    xfun_0.26
## [17] utf8_1.2.2      withr_2.4.2    htmltools_0.5.2 ellipsis_0.3.2
## [21] yaml_2.2.1      digest_0.6.28  tibble_3.1.5    lifecycle_1.0.1
## [25] crayon_1.4.1    farver_2.1.0   purrr_0.3.4     vctrs_0.3.8
## [29] glue_1.4.2      evaluate_0.14  rmarkdown_2.11  labeling_0.4.2
## [33] stringi_1.7.5   compiler_4.1.1 pillar_1.6.3    generics_0.1.0
## [37] scales_1.1.1    pkgconfig_2.0.3
```